

```
# Step 1 - Import all required libraries
```

```
import pandas as pd
from sklearn.preprocessing import StandardScaler
import matplotlib.pyplot as plt
from sklearn.cluster import KMeans
```

```
# Step 2 - Load the dataset (sales_data_sample.csv)
```

```
df = pd.read_csv("sales_data_sample.csv", encoding ='latin1')
df.head()
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	ORDERDATE	STATUS	QTR_ID	MONTH_ID	YEAR_ID	...	ADDRESSLINE1	ADRE
0	10107	30	95.70	2	2871.00	2/24/2003 0:00	Shipped	1	2	2003	...	897 Long Airport Avenue	
1	10121	34	81.35	5	2765.90	5/7/2003 0:00	Shipped	2	5	2003	...	59 rue de l'Abbaye	
2	10134	41	94.74	2	3884.34	7/1/2003 0:00	Shipped	3	7	2003	...	27 rue du Colonel Pierre Avia	
3	10145	45	83.26	6	3746.70	8/25/2003 0:00	Shipped	3	8	2003	...	78934 Hillside Dr.	
4	10159	49	100.00	14	5205.27	10/10/2003 0:00	Shipped	4	10	2003	...	7734 Strong St.	

5 rows × 25 columns

```
# Step 3 - Select only numeric columns for clustering
```

```
X = df.select_dtypes(include=['float64', 'int64']).dropna()
X.head()
```

	ORDERNUMBER	QUANTITYORDERED	PRICEEACH	ORDERLINENUMBER	SALES	QTR_ID	MONTH_ID	YEAR_ID	MSRP	grid icon	info icon
0	10107	30	95.70	2	2871.00	1	2	2003	95		
1	10121	34	81.35	5	2765.90	2	5	2003	95		
2	10134	41	94.74	2	3884.34	3	7	2003	95		
3	10145	45	83.26	6	3746.70	3	8	2003	95		
4	10159	49	100.00	14	5205.27	4	10	2003	95		

Next steps: [Generate code with X](#) [New interactive sheet](#)

💡 What this does (simple words):

Code	Meaning
<code>(select_dtypes(include=['float64', 'int64']))</code>	Keeps only numeric columns (like SALES, QUANTITYORDERED, etc.).
<code>.dropna()</code>	Removes rows that have missing (NaN) values — keeps data clean.
<code>X.head()</code>	Shows first 5 rows of the numeric-only data.

```
# Step 4 - Scale (Normalize) the data
```

```
scaler = StandardScaler()
X_scaled = scaler.fit_transform(X)
```

💡 What this does in simple words:

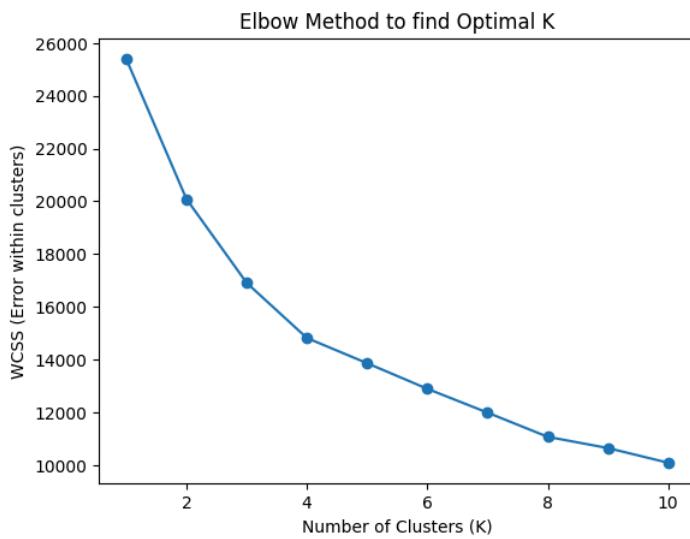
Code	Meaning
<code>StandardScaler()</code>	Selects method to scale/normalize data.
<code>fit_transform(X)</code>	Learns the scaling from the data and applies it.
<code>X_scaled</code>	Now contains the same data but standardized → perfect for K-Means.

```
# Step 5 - Use the Elbow Method to find the best number of clusters (K)
```

```
wcss = [] # WCSS = Within-Cluster Sum of Squares (measures how tight the clusters are)
```

```
for i in range(1, 11): # Try K = 1 to 10 clusters
    kmeans = KMeans(n_clusters=i, random_state=42)
    kmeans.fit(X_scaled)
    wcss.append(kmeans.inertia_) # inertia_ = WCSS for that K
```

```
# Plot the Elbow Graph
plt.plot(range(1, 11), wcss, marker='o')
plt.title('Elbow Method to find Optimal K')
plt.xlabel('Number of Clusters (K)')
plt.ylabel('WCSS (Error within clusters)')
plt.show()
```



💡 What this does:

Tries different values of K (1 to 10).

Calculates WCSS (how spread out the points are inside each cluster).

Plots a graph:

X-axis → Number of clusters (K)

Y-axis → WCSS

We look for a point that looks like an elbow/bend in the graph → that's the best K.

Eg. If the graph sharply drops until K=3 and then slows → K = 3 is best.

```
# Step 6 – Apply K-Means clustering using the best K value
```

```
kmeans = KMeans(n_clusters = 3, random_state = 42)
df['Cluster'] = kmeans.fit_predict(X_scaled)
```

💡 What this means (super simple):

Code	Meaning
KMeans(n_clusters=3)	We are telling the model to form 3 groups/clusters.
fit_predict(X_scaled)	Groups the data and returns which cluster each row belongs to.
df['Cluster']	Adds a new column named Cluster to the original dataset to store results.

```
# Step 7 – View the clustering results
```

```
# Count how many rows are in each cluster
print(df['Cluster'].value_counts())

# Show first few rows with cluster labels
print("\nSample data with Cluster labels:\n")
print(df[['ORDERNUMBER', 'SALES', 'Cluster']].head())
```

```
Cluster
1    1183
2     907
0     733
Name: count, dtype: int64
```

Sample data with Cluster labels:

```
ORDERNUMBER    SALES  Cluster
0        10107  2871.00      1
1        10121  2765.90      2
2        10134  3884.34      1
3        10145  3746.70      1
4        10159  5205.27      1
```

💡 What this shows:

Output	Meaning
value_counts()	Tells how many records belong to Cluster 0, 1, 2, etc.
df[['ORDERNUMBER', 'SALES', 'Cluster']]	Shows ORDERNUMBER, its SALES value, and which cluster it belongs to.

